

TiDB Prometheus 日志增长临时解决方案

问题原因:

该问题是由于该版本中监控指标 `tikv_tiflash_proxy_thread_nonvoluntary_context_switches` 以及其他类似的指标按照 `thread_name` 和 `thread_id` 进行分类, 但是 `tiflash` 的执行模型会创建较多的 `thread`, 导致该监控指标产生了大量的 `label` 并且在每次上报时都会持续上报这些指标, 因此导致监控数据上涨较快, 重启 `tiflash` 可以解决该问题。另外该问题已经在 `v4.0.9+` 的版本中修复, 可以考虑升级到修复后的版本。

Workaround:

在集群 `prometheus.yml` 中 `scrape_configs.job_name: "tiflash"` 添加如下配置可以屏蔽造成该问题的指标。

```
metric_relabel_configs:
- source_labels: [__name__]
  regex:
'tikv_tiflash_proxy_thread_nonvoluntary_context_switches|tikv_tiflash_proxy_thread_voluntary_context_switches|tikv_tiflash_proxy_thread_cpu_seconds_total|tikv_tiflash_proxy_threads_io_bytes_total'
  action: drop
```

```
vim /data1/tidb/deploy/prometheus-9090/conf/prometheus.yml
```

```
- job_name: "tiflash"

  metric_relabel_configs:

    - source_labels: [__name__]

      regex:
'tikv_tiflash_proxy_thread_nonvoluntary_context_switches|tikv_tiflash_proxy_thread_voluntary_context_switches|tikv_tiflash_proxy_thread_cpu_seconds_total|tikv_tiflash_proxy_threads_io_bytes_total'

      action: drop


- job_name: "overwritten-nodes"

  metric_relabel_configs:

    - source_labels: [__name__]

      regex: 'node_interrupts_total'

      action: drop
```

注意: `tiup` 不支持变更 `prometheus` 配置, 需要手动编辑对应组件的配置文件, 然后 `kill -1 prometheus` 进程, 等待其自动重启后即可生效。另外以后对集群进行的重启, 扩缩容等操作都会覆盖修改后的配置文件, 所以需要重复该步骤。

```
- job_name: "tiflash"
  honor_labels: true # don't overwrite job & instance labels
  static_configs:
    - targets:
      - '172.16.5.59:8344'
      - '172.16.5.59:20592'

  metric_relabel_configs:
    - source_labels: [__name__]
      regex: 'tikv_tiflash_proxy_thread_nonvoluntary_context_switches|tikv_tiflash_proxy_thread_voluntary_context_switches|tikv_tiflash_proxy_thread_cpu_seconds_total'
      action: drop
```

```
scrape_configs:
- job_name: "overwritten-nodes"
  honor_labels: true # don't overwrite job & instance labels
  static_configs:
  - targets:
    - '100.84.98.30:9100'
    - '100.84.98.17:9100'
    - '100.84.97.163:9100'
    - '100.84.98.31:9100'
    - '100.84.98.32:9100'
    - '100.84.97.8:9100'
    - '100.84.98.42:9100'
    - '100.84.98.43:9100'
    - '100.84.98.44:9100'
    - '100.84.98.18:9100'
    - '100.84.98.15:9100'
    - '100.84.98.16:9100'
  metric_relabel_configs:
  - source_labels: [__name__]
    regex: 'node_interrupts_total'
    action: drop
job_name: "c1db"
honor_labels: true # don't overwrite job & instance labels
```